



# Performance analysis on GPUs with NVIDIA tools

Dominik Ernst



# Example: 2D Jacobi

Log in to Alex:

```
ssh -J <user>@csnhr.nhr.fau.de <user>@alex
```

Allocate an interactive Job on a GPU:

```
srun --gres=gpu:a100:1 -p a100 --reservation  
GPUperf --time 4:00:00 --pty /bin/bash -l
```

Allocate an interactive Job from different account:

```
srun --gres=gpu:a100:1 -p a100 -C a100_40 --time  
4:00:00 --pty /bin/bash -l
```

# Example: 2D Jacobi

Get the Source:

```
git clone https://github.com/te42kyfo/omp_jacobi  
or
```

```
cp -r ../b53k0000/omp_jacobi
```

Load the compiler module

```
module load nvhpc  
module load cuda
```

Build and run the CPU base line

```
make main1  
./main1
```

# Roofline Analysis: 2D Jacobi

3 ADDs, 1 MUL per iteration:

$$A[o] = 0.25 * (B[^] + B[v] + B[<] + B[>])$$

Read entire grid A and B once each → on average: one value per iteration

$$= 2 \times 8B / \text{iteration}$$

Code Intensity:

$$4 \text{ Flop} / 16 B = 0.25 \text{ Flop/B}$$

A100 Machine Intensity:

$$9.7 \text{ Tflop/s} / 1555 \text{ GB/s} = 6.2 \text{ Flop/B}$$



# Build/Run/Profile

Build and run the Nth version

```
make main<N>  
./main<N>
```

Create a profile

```
nsys profile main<N>
```

Launch the profiling GUI

```
nsys-ui
```

# Kernel Profiling

## Kernel profiling

```
ncu <application>
```

## List metric sections

```
ncu --list-sections
```

## Collect all sections

```
ncu --set full -f -o <output file> <application>
```

## Launch the ncu profiling GUI

```
ncu-ui
```

# GPU Architecture

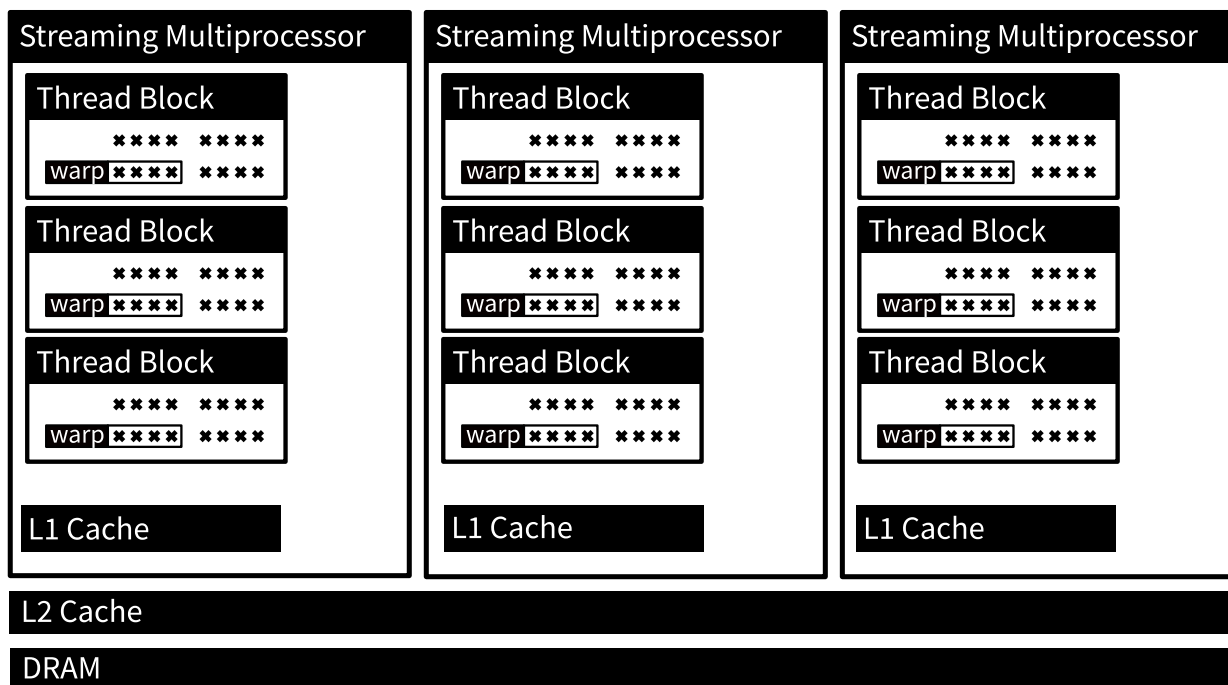
32 threads → 1 warp

up to 1024 threads / 32 warps → 1 thread block


up to 64 warps / 2048 threads → 1 SM

108 SM → A100 GPU

2048 threads / SM \* 108SM → ~200'000 threads / GPU



# GPU Architecture



per SM: 192 kB L1 cache  
shared for all SM: 40MB L2 cache  
shared for all SM: 40 GB DRAM

(A100-SXM4-40GB)




# OpenMP Loop main4

```
#pragma omp target parallel for
for (int y = 1; y < height - 1; y++)
    for (int x = 1; x < width - 1; x++)
        A[y][x] = ...
```




```
kernel:
    y = global_thread_id.x;
    for (int x = 1;
         x < width - 1;
         x++)
        A[y][x] = ...
```

```
kernel<<< rows >>>(...)
```



# OpenMP Loop main41

```
#pragma omp target parallel for  
for (int x = 1; x < width - 1; x++)  
    for (int y = 1; y < height - 1; y++)  
        A[y][x] = ...
```



# OpenMP main51


```
#pragma omp target parallel for collapse(2)
for (int y = 1; y < height - 1; y++)
    for (int x = 1; x < width - 1; x++)
        A[y][x] = ...
```



kernel:

```
y = global_thread_id.x / N;
X = global_thread_id.x % N;
A[y][x] = ...
```

```
kernel<<< rows*columns >>>(...)
```



# OpenMP Loop main6 / main7

```
#pragma omp target parallel for collapse(2)
for (int oy = 1; oy < height - 1; oy += 4)
    for (int x = 1; x < width - 1; x++)
        for (int iy = 0; iy < 4; iy++) {
            int y = oy + iy;
```

